

Applied Programming

Arrays and Pointers

in C

More details in: G. Semeraro, "Numerical Programming in C", chapter 4, (in Mycourses)

Arrays and Pointers

- In C, *arrays of primitive types* are implemented as blocks of contiguous memory.

```
int a[10];          /* 1D array of 10 integers */
double A[3][5];     /* 2D array of floats ("3 rows, 5 columns") */
```

- Arrays are can be initialized as follows

```
int a[] = {1,2,6,7,8}; /* 1D array of 5 integers */
```

- Array elements can be accesses by *indexing* and **pointer arithmetic**

Ex: The following are all equivalent `A[i][j]`

`*(A[i] + j)`

`*((A+i))[j]`

`*((*(A+i)) + j)`

`*(&A[0][0] + 5*i+j)`

The compiler uses “a storage mapping function” to convert the indices to a memory address.

Arrays and Pointers

- In C, the **name of an array** is a **pointer** to the memory address where the first array value is stored

a same as **&a[0]** **/* 1D array */**

A same as **&A[0][0]** **/* 2D array */**

- In our previous example, the *storage mapping function* for the 2D array **A[3][5]** is

$$s(i, j) = 5*i + j$$

- Notes:
 - The storage mapping function is used by the compiler to generate object code to access the correct array element in memory. The compiler *only needs* to know the *last dimension of the array* (e.g., **5**) to generate the *correct memory offset*.
 - The compiler uses pointer arithmetic.
If **ptr=&A[0][0]** then **A[2][3]** is ***(ptr + 5*(2)+(3))**

Memory Storage of 2D Arrays

- C stores 2D-arrays in *row-major form*, (as in Java and C++). Some languages, e.g., FORTRAN, use column-major form.

Example: The 2D array `A[3][3]` in memory.

2D array	Memory storage																		
<table><tr><td>1</td><td>2</td><td>3</td></tr><tr><td>4</td><td>5</td><td>6</td></tr><tr><td>7</td><td>8</td><td>9</td></tr></table>	1	2	3	4	5	6	7	8	9	<table><tr><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td></tr></table>	1	2	3	4	5	6	7	8	9
1	2	3																	
4	5	6																	
7	8	9																	
1	2	3	4	5	6	7	8	9											

Warning: When calling FORTRAN code (e.g., optimized linear algebra routines) from C code this must be taken into account !

Using Arrays


- The most common use of arrays in numerical computations is to **represent Matrices**.
- A straight forward representation of a “matrix” in C is using a “fixed” 2D array

```
#define N_ROWS (20)  
#define N_COLS (40)  
double C[N_ROWS][N_COLS];
```

- Practical Limitation
Cannot use 2D arrays in function calls unless we know **dimensions *a priori*** (and in practice we almost never do !)

How to we address this limitation ?

Dynamically Allocated Arrays



Dynamically Allocated 2D Array in
Memory ("row-major" form)

Dynamically Allocated Arrays

- A better way to represent a matrix in C is as pointers to (arrays of) pointers
- The size of the matrix neither needs to be known a priori nor be coded as a constant.

```
#include <stdlib.h>      /* for calloc and malloc */
typedef double MatElement;

. . .
int i, /* row index counter */
nr,nc; /* # rows and # columns */

MatElement **A = NULL; /* declare matrix A */
. . .
/* Allocate space for array of pointers */
A = malloc( nr * sizeof(MatElement *));
/* Allocate space for matrix elements and set to zero */
for (i=0; i<nr; i++)
    A[i] = calloc( nc, sizeof(MatElement) );
```

Representing (Dense) Matrices

- The space *must be* allocated in **row-major form** to be able to access its elements using the *standard indexing syntax* **A[i][j]**

```
#include <stdlib.h>    /* for malloc      */
typedef double MatElement;
int      i,j,nr,rc;
MatElement **A = NULL;    /* declare matrix */
. . . .
A = malloc( nr*sizeof(MatElement *) );
for (i=0; i<nr; i++)
    A[i] = malloc( nc*sizeof(MatElement) );
. . . .
for (i=0; i<nr; i++) /* Initialize to random */
    for (j=0; j<nc; j++)
        A[i][j] = rand();
```


Releasing Matrix Space

- To release the allocated memory space “*undo*” the allocation operations in reverse order.

```
void free_matrix_space(MatElement **A, int m) {  
    int i;                                /* row counter */  
    for (i=0; i<m; i++)  
        free(A[i]);    /* free rows */  
    A[i] = NULL;    /* “grounding” pointers */  
    free(A); /* Then column vector of row pointers */  
    A = NULL;    /* grounding */  
}
```

Warning: Doing `free A` will not work !

Efficient Matrix Allocation

- It is *more efficient* to *allocate all memory space at once* as shown below

```
#include <stdlib.h>      /* for calloc and malloc */
typedef double MatElement;

. . .
int i,                  /* row index counter      */
nr,nc;                 /* # rows and # columns */
MatElement *ptr = NULL; /* pointer to rows      */
MatElement **A = NULL; /* ptr to matrix        */

. . .
/* Allocate array of row pointers */
A = malloc(nr * sizeof(MatElement *));
/* Allocate all matrix at once */
ptr = calloc(nr * nc, sizeof(MatElement));
/* Set array of row pointers properly */
for (i=0;i<nr;i++)
    A[i] = ptr + nc*i;
```

Efficient Matrix Allocation

- You can release the memory of an array allocated all at once as follows:

```
void free_all_matrix (MatElement **A) {  
    free(A[0]); /* free array of elements */  
    A[0] = NULL;  
    free(A);    /* free array of ptrs to rows */  
    A = NULL;  
}
```

Note: We do not need to pass the size of the array for deallocation (this is not the case if we allocate space row by row).

Example: row_major.c

- Allocate using row by row: A [6][3]
- Allocate all at once: C [6][3]

Elements of A (row by row allocation)

1	(10000)	2	(10004)	3	(10008)
4	(10032)	5	(10036)	6	(10040)
7	(10064)	8	(10068)	9	(10072)
10	(10096)	11	(10100)	12	(10104)
13	(10128)	14	(10132)	15	(10136)
16	(10160)	17	(10164)	18	(10168)

Total bytes: 172

Elements of C (all at once allocation)

1	(10000)	2	(10004)	3	(10008)
4	(10012)	5	(10016)	6	(10020)
7	(10024)	8	(10028)	9	(10032)
10	(10036)	11	(10040)	12	(10044)
13	(10048)	14	(10052)	15	(10056)
16	(10060)	17	(10064)	18	(10068)

Total bytes: 72

(Memory address in parenthesis)

Why does the row by row allocation take more memory?

Example: Changing Indexing

- C is a very flexible programming language so it allows you to do many things (you are in charge)
- In most *matrix computations algorithms indexing starts at 1*, e.g., $A[1][1]$ is the first element of matrix A. That is also how it is in **FORTAN** (**FOR**mula **TRAN**slator)
- Here we show an unorthodox way to “change indexing” of matrices so that it starts at 1 (instead of 0 as is the default in C 2d arrays)

Adjusting Subscript Range

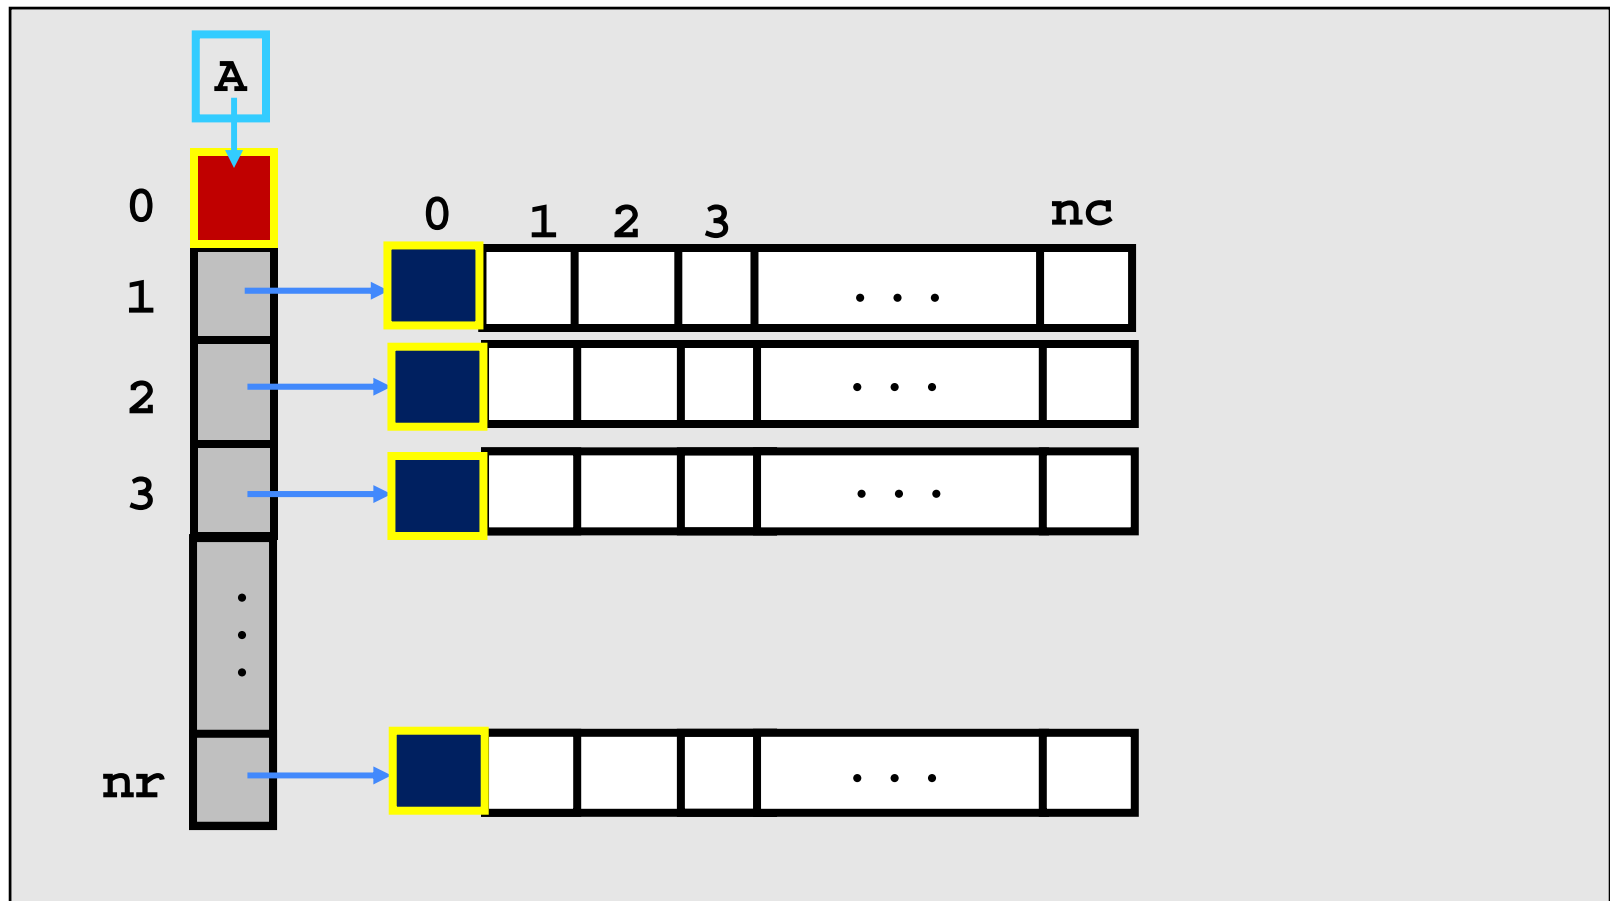
- Most mathematical algorithms use subscripts that start at 1 not at 0

Q: How do we handle that in C ?

A: Offset the pointers

```
typedef double MatElement;
MatElement **matrix1_simpleAlloc(int m, int n) {
    int i;
    MatElement **A = NULL;
    A = malloc(m * sizeof(MatElement*)); /*m x n matrix */
    --A; /* offset row pointer back by one */
    for (i=1; i<=m; ++i) {
        A[i] = calloc( n, sizeof(MatElement));
        --A[i]; /* offset col pointer back by one */
    } /* End for i*/
    return(A); } /* End matrix1_simpleAlloc */
```

Dense Matrix: Dynamically Allocated



Dynamically Allocated Matrix with offset

Adjusting Subscript Range

- The following sample code illustrates the use of *subscript adjusted matrices*

```
typedef double MatElement;  
MatElement **A;  
....  
  
A = matrix1_simpleAlloc(m,n); /* offset mxn matrix */  
for (i=1; i<=m; i++)  
    for (j=1; j<=n; j++)  
        A[i][j] = ....  
  
....  
matrix1_simpleFree(A, m, n) /* how ? */
```

- How do we *de-allocate* the memory space ?

Releasing Matrix Space

- To release the allocated memory space we need to point to the right location

```
void matrix1_simpleFree(MatElement **A, int m, int n)
{
    int i;
    for (i=1; i<=m; ++i)
    {
        /* Return the col pointer BACK to the
           original then FREE */
        free(++A[i]);
        A[i] = NULL;
    }
    /* Return the row pointer BACK to the
       original, THEN free */
    free(++A);
    A = NULL;
} /* End matrix1_simmpleFree */
```

More efficient version

- The following version with allocation of matrix space all at once is more efficient

```
MatElement **matrix1_cleverAlloc(int m, int n) {
    int i;                      /* row index counter */
    MatElement *ptr=NULL;       /* pointer to rows */
    MatElement **A=NULL;        /* ptr to matrix */

    /* Allocate array of row pointers and shift back */
    A = malloc(m * sizeof(MatElement *));
    --A;
    /* Allocate all matrix data at once */
    ptr = calloc( m * n, sizeof(MatElement));
    /* Set array of row pointers properly */
    for (i = 1; i <= m; ++i)
    {
        A[i] = ptr + (n*(i-1)) - 1;
    }
    return(A);
}
```

More efficient Version

- To release the allocated memory space need to point to the right location

```
void matrix1_cleverFree (MatElement **A) {  
    free(A[1]+1); /* base address of array elements */  
    A[1] = Null;  
    free(A+1);  
    A = NULL;  
}
```

- Note that, in this case, deallocation is simple, we don't need the dimensions of A.

Summary: Arrays and Matrices in C

- *Dense Matrices* are usually declared as *arrays of pointers to pointers* (to the element type of interest) and allocated dynamically.
- **Vectors** (1D arrays) are naturally represented as a *pointers* (to the element type of interest) and allocated dynamically as well.
- Regular *matrix indexing* can be used to access the *elements* of a matrix.
- *Matrices in C are stored in row-major form.* Therefore to use indexing we need to dynamically allocate space also in “row-major form”

Strings

- Strings: (not a data type)
 - **Array of characters** terminated by the **sentinel \0**
 - Constant strings enclosed with double quotes:
 - The string : **"A" [65][\0]**
 - The character : **'A' [65]**

Examples:

```
/* declare and initialize a string */  
char[ ] course="CMPE-380";  
course[2]=='P'; /* is True */
```

- The standard library **<string.h>** has functions to operate on strings, such as:
strcpy, strcmp, strcat, strstr, strchr

Comparing Strings

- **You can't compare strings directly in C**
 - Only character by character
- Given:

```
char *str1 = "Some Text1";  
char *str2 = "Some Text2";
```
- Consider: **if (str1 == str2)**
 - **NEVER** true, only comparing **POINTERS** str1 and str2
- Consider: **if (*str1 == *str2)**
 - **ALWAYS** true in this case
 - Only comparing the **FIRST CHARACTERS** in str1 and str2
- Use the function **strcmp()**

strcmp()

```
int strcmp(const char *s1, const char *s2);
```

- s1, s2 - *strings to be compared, byte by byte*
- Returns –
 - 0 - if strings are identical
 - Negative - if, character by character, s1 is less than s2
 - Positive - if, character by character. s1 is greater than s2

How it works:

strcmp("abc", "ABC") is *positive*

because “a” (0x61) is greater than “A” (0x41)

Use: if (**!strcmp(s1, s2)**) { printf(**“Identical”**);}

strcpy()

```
char *strcpy(char *dest, const char *src);
```

- Copies **src** into **dest**
- byte by byte **until 0x00** in src is reached
- *****Assumes***** dest is big enough (dangerous)
- Returns – Pointer to dest

How it works:

```
char dest [20];
```

```
strcpy(dest, "ABC");
```

dest [0] contains “A”

dest [1] contains “B”

dest [2] contains “C”

dest [3] contains 0x00

strncpy()

```
char *strncpy(char *dest, const char *src , int num);
```

- Copies **src** into **dest**
- byte by byte **until 0x00** in src is reached
- Will not copy more than **num** bytes (**safer**)
 - *Won't overflow the destination buffer!*
- Returns – Pointer to dest

How it works:

```
char dest [20];  
strncpy(dest, "ABC", 20);
```

```
dest [0] contains "A"  
dest [1] contains "B"  
dest [2] contains "C"  
dest [3] contains 0x00
```

strcat()

```
char *strcat(char *dest, const char *src);
```

- Appends **src** to the **END of dest**
- byte by byte until **0x00** is src is reached
- Assumes dest is big enough
- Returns – Pointer to dest

How it works:

```
char dest [20];
```

```
strcpy(dest, "abc");
```

```
strcat(dest, "123");
```

results in dest containing "abc123"

strncat()

```
char *strncat(char *dest, const char *src, int num);
```

- Appends **src** to the **END of dest**
- byte by byte until **0x00** is src is reached
- Will not copy more than **num** bytes (**safer**)
- Returns – Pointer to dest

How it works:

```
char dest [20];
```

```
strcpy(dest, "abc", 20);
```

```
strcat(dest, "123", 20);
```

results in dest containing "abc123"

Problem 1

Crashes when I use “poly” in main.

What's **wrong?**

main.c:

```
polynomial poly;
```

```
createPoly(poly, 10);
```

```
poly.polyCoef[0] = 1;      /* crash */
```

```
.....
```

```
void createPoly(polynomial *p, int n){
```

```
    p = malloc(sizeof(polynomial));
```

```
    p->polyCoef = malloc(sizeof(complex) * n);
```

```
}
```

*Error checking
removed for
clarity*

Problem 1

- Addresses for example only

polynomial **poly**

```
createPoly(poly, 10);  
createPoly(polynomial  
*p, int n)
```

Creates storage **@100**

passes the value p= 100, n= 10

```
p = malloc();
```

```
p->polyCoef = malloc()
```

A new p is created **@400**

400->polyCoef = **@1000**

Back in main

```
poly.polyCoef [0] = 1
```

100.polyCoef undefined!

Problem 2

Crashes when I use “poly” in main.

main.c:

```
polynomial *poly;
```

```
poly = malloc(sizeof(polynomial));
```

```
createPoly(poly, 10);
```

```
poly.polyCoef[0] = 1;      /* crash */
```

.....

```
void createPoly(polynomial *p, int n){
```

```
    p = malloc(sizeof(polynomial));
```

```
    p->polyCoef = malloc(sizeof(complex) * n);
```

```
}
```

Just like before, changing the pointer
in the subroutine does nothing!

Question

- What would **strcat("123", "456")** do?
- Logically bad things.
 - The compiler puts "123" **someplace in memory**
 - strcat will try to **append "456" after the end of "123"** BUT
 - the compiler has put **other variables in ram** after the "123" SO
 - this operation will result in those other variables **being destroyed**